

# CS 33

## Files Part 3

# The File Abstraction

- A file is a simple array of bytes
- A file is made larger by writing beyond its current end
- Files are named by paths in a naming tree
- System calls on files are synchronous
- Files are permanent

# Naming

- **(almost) everything has a path name**
  - files
  - directories
  - devices (known as *special files*)
    - » keyboards
    - » displays
    - » disks
    - » etc.

# I/O System Calls

- **int** file\_descriptor = open(pathname, mode [, permissions])
- **int** close(file\_descriptor)
- **ssize\_t** count = read(file\_descriptor, buffer\_address, buffer\_size)
- **ssize\_t** count = write(file\_descriptor, buffer\_address, buffer\_size)
- **off\_t** position = lseek(file\_descriptor, offset, whence)

# Standard File Descriptors

```
int main( ) {  
    char buf[BUFSIZE];  
    int n;  
    const char *note = "Write failed\n";  
  
    while ((n = read(0, buf, sizeof(buf))) > 0)  
        if (write(1, buf, n) != n) {  
            write(2, note, strlen(note));  
            exit(1);  
        }  
    return(0);  
}
```

# Standard I/O Library

Formatting

**printf** ... **scanf**

Buffering

**stdin** **stdout** **stderr** ...

Syscalls

**fd 0** **fd 1** **fd 2** ...

# Standard I/O

```
FILE *stdin;           // declared in stdio.h
FILE *stdout;          // declared in stdio.h
FILE *stderr;          // declared in stdio.h
```

```
scanf("%d", &in);      // read via f.d. 0
printf("%d\n", in);    // write via f.d. 1
fprintf(stderr, "there was an error\n");
    // write via f.d. 2
```





# Unbuffered Output

```
fprintf(stderr, "xy");  
fprintf(stderr, "zz");  
fprintf(stderr, "y\n");
```

**x y z z y**

**display**

# I/O System Calls

- **int** file\_descriptor = open(pathname, mode [, permissions])
- **int** close(file\_descriptor)
- **ssize\_t** count = read(file\_descriptor, buffer\_address, buffer\_size)
- **ssize\_t** count = write(file\_descriptor, buffer\_address, buffer\_size)
- **off\_t** position = lseek(file\_descriptor, offset, whence)

# Standard File Descriptors

```
int main( ) {  
    char buf[BUFSIZE];  
    int n;  
    const char *note = "Write failed\n";  
  
    while ((n = read(0, buf, sizeof(buf))) > 0)  
        if (write(1, buf, n) != n) {  
            write(2, note, strlen(note));  
            exit(1);  
        }  
    return (0);  
}
```

# Standard I/O Library

Formatting

**printf** ... **scanf**

Buffering

**stdin** **stdout** **stderr** ...

Syscalls

**fd 0** **fd 1** **fd 2** ...

# Standard I/O

```
FILE *stdin;           // declared in stdio.h
FILE *stdout;          // declared in stdio.h
FILE *stderr;          // declared in stdio.h

scanf("%d", &in);       // read via f.d. 0
printf("%d\n", in);     // write via f.d. 1
fprintf(stderr, "there was an error\n");
                        // write via f.d. 2
```



# Unbuffered Output

```
fprintf(stderr, "xy");  
fprintf(stderr, "zz");  
fprintf(stderr, "y\n");
```

**x y z z y**

**display**

# A Program

```
int main(int argc, char *argv[]) {  
    if (argc != 2) {  
        fprintf(stderr, "Usage: echon reps\n");  
        exit(1);  
    }  
    int reps = atoi(argv[1]);  
    if (reps > 2) {  
        fprintf(stderr, "reps too large, reduced to 2\n");  
        reps = 2;  
    }  
    char buf[256];  
    while (fgets(buf, 256, stdin) != NULL)  
        for (int i=0; i<reps; i++)  
            fputs(buf, stdout);  
    return (0);  
}
```



# From the Shell ...

```
$ echon 1
```

- ***stdout*** (fd 1) and ***stderr*** (fd 2) go to the display
- ***stdin*** (fd 0) comes from the keyboard

```
$ echon 1 > Output
```

- ***stdout*** goes to the file “Output” in the current directory
- ***stderr*** goes to the display
- ***stdin*** comes from the keyboard


```
$ echon 1 < Input
```

- ***stdin*** comes from the file “Input” in the current directory


# Redirecting Stdout in C

```
if ((pid = fork()) == 0) {  
    /* set up file descriptor 1 in the child process */  
    close(1);  
    if (open("/home/twd/Output", O_WRONLY) == -1) {  
        perror("/home/twd/Output");  
        exit(1);  
    }  
    char *argv[] = {"echon", "2", NULL};  
    execv("/home/twd/bin/echon", argv);  
    exit(1);  
}  
  
/* parent continues here */  
  
waitpid(pid, 0, 0);    // wait for child to terminate
```


# File-Descriptor Table



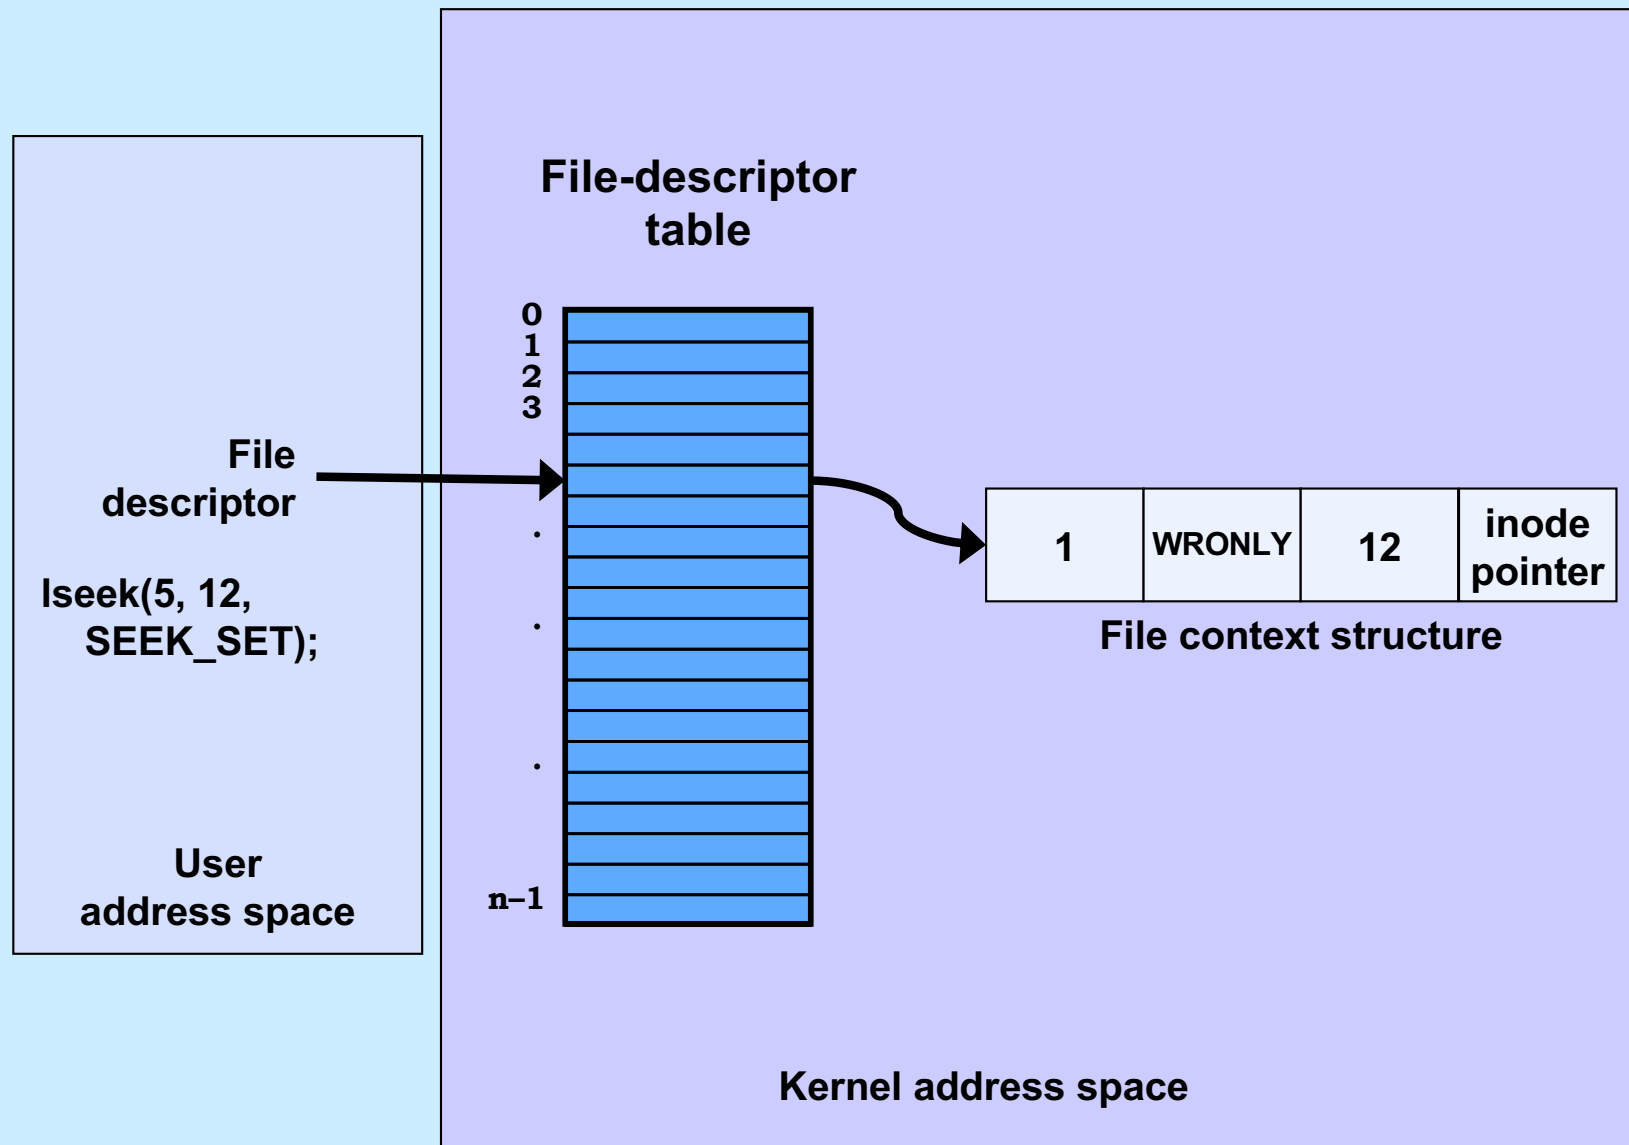
# File Location



# File Location



# File Location



# Allocation of File Descriptors

- Whenever a process requests a new file descriptor, the lowest-numbered file descriptor not already associated with an open file is selected; thus

```
#include <fcntl.h>
#include <unistd.h>
```

```
close(0);
fd = open("file", O_RDONLY);
```

- will always associate *file* with file descriptor 0 (assuming that *open* succeeds)

# Redirecting Output ... Twice

```
if (fork() == 0) {  
    /* set up file descriptors 1 and 2 in the child process */  
    close(1);  
    close(2);  
    if (open("/home/twd/Output", O_WRONLY) == -1) {  
        exit(1);  
    }  
    if (open("/home/twd/Output", O_WRONLY) == -1) {  
        exit(1);  
    }  
    char *argv[] = {"echon", 2, NULL};  
    execv("/home/twd/bin/echon", argv);  
    exit(1);  
}  
/* parent continues here */
```

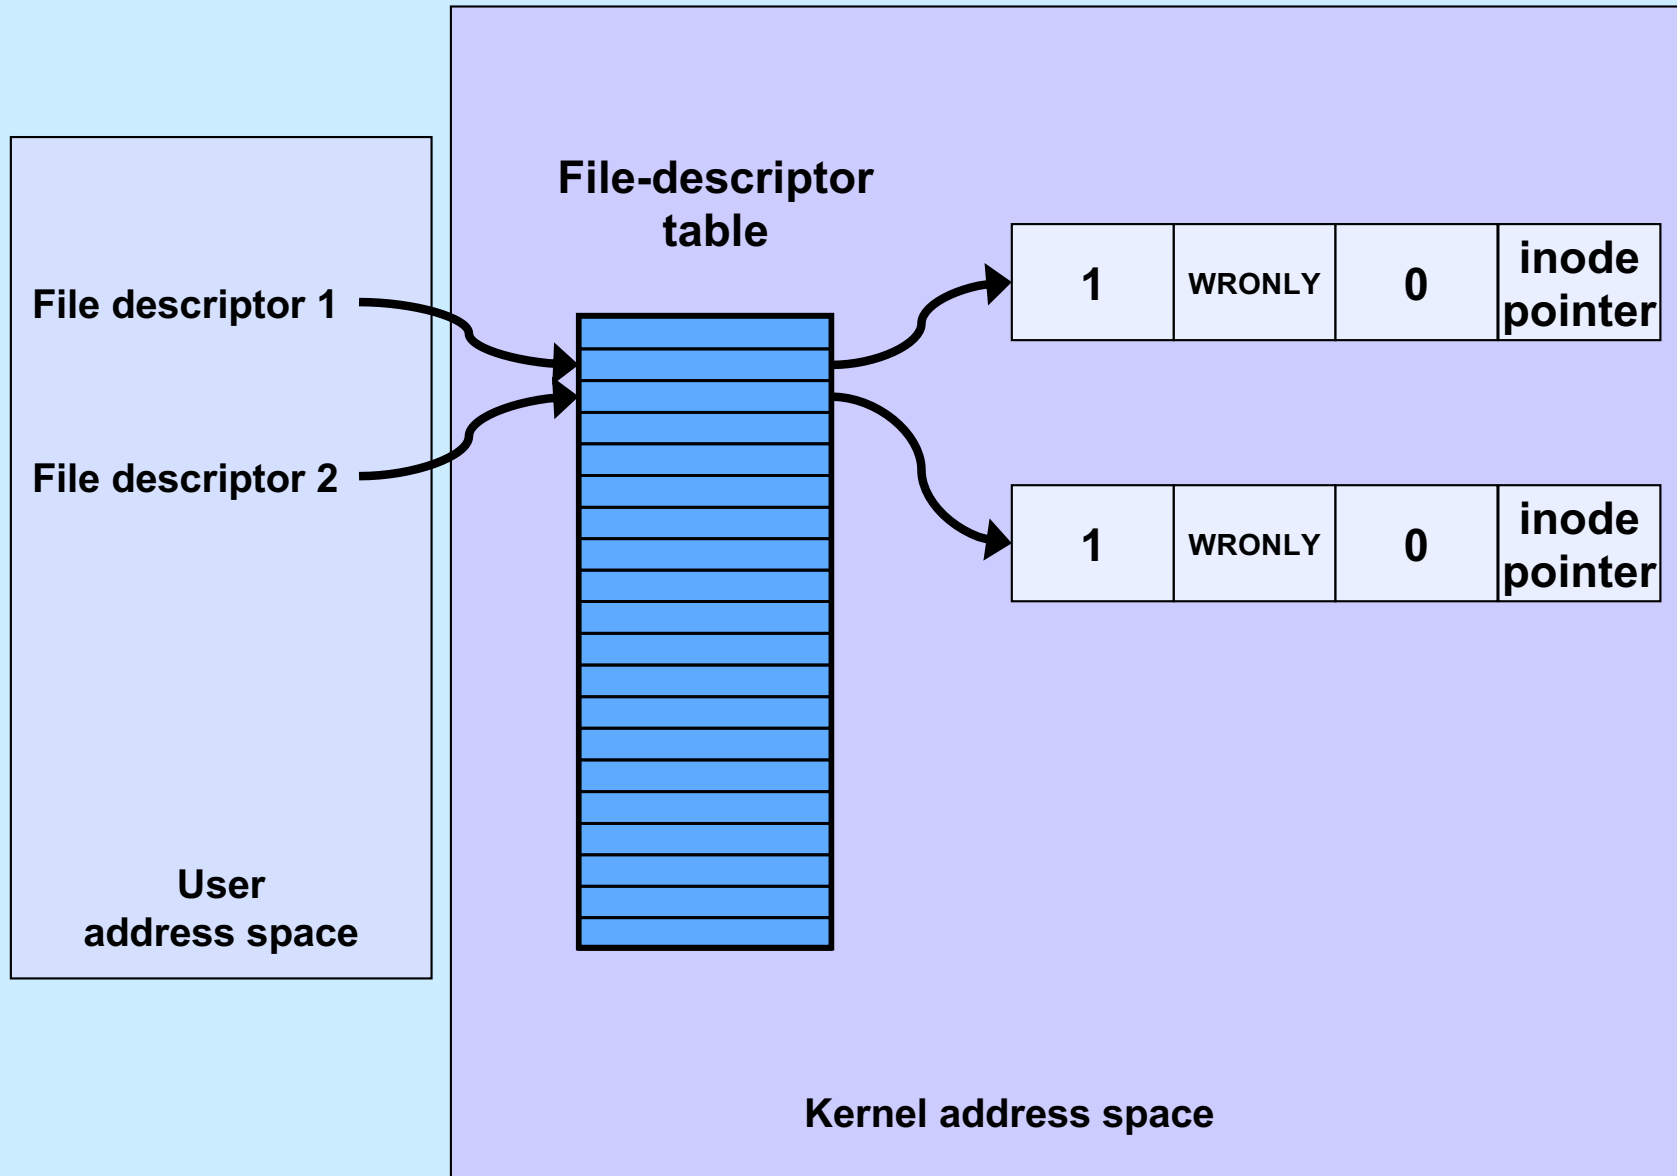


# From the Shell ...

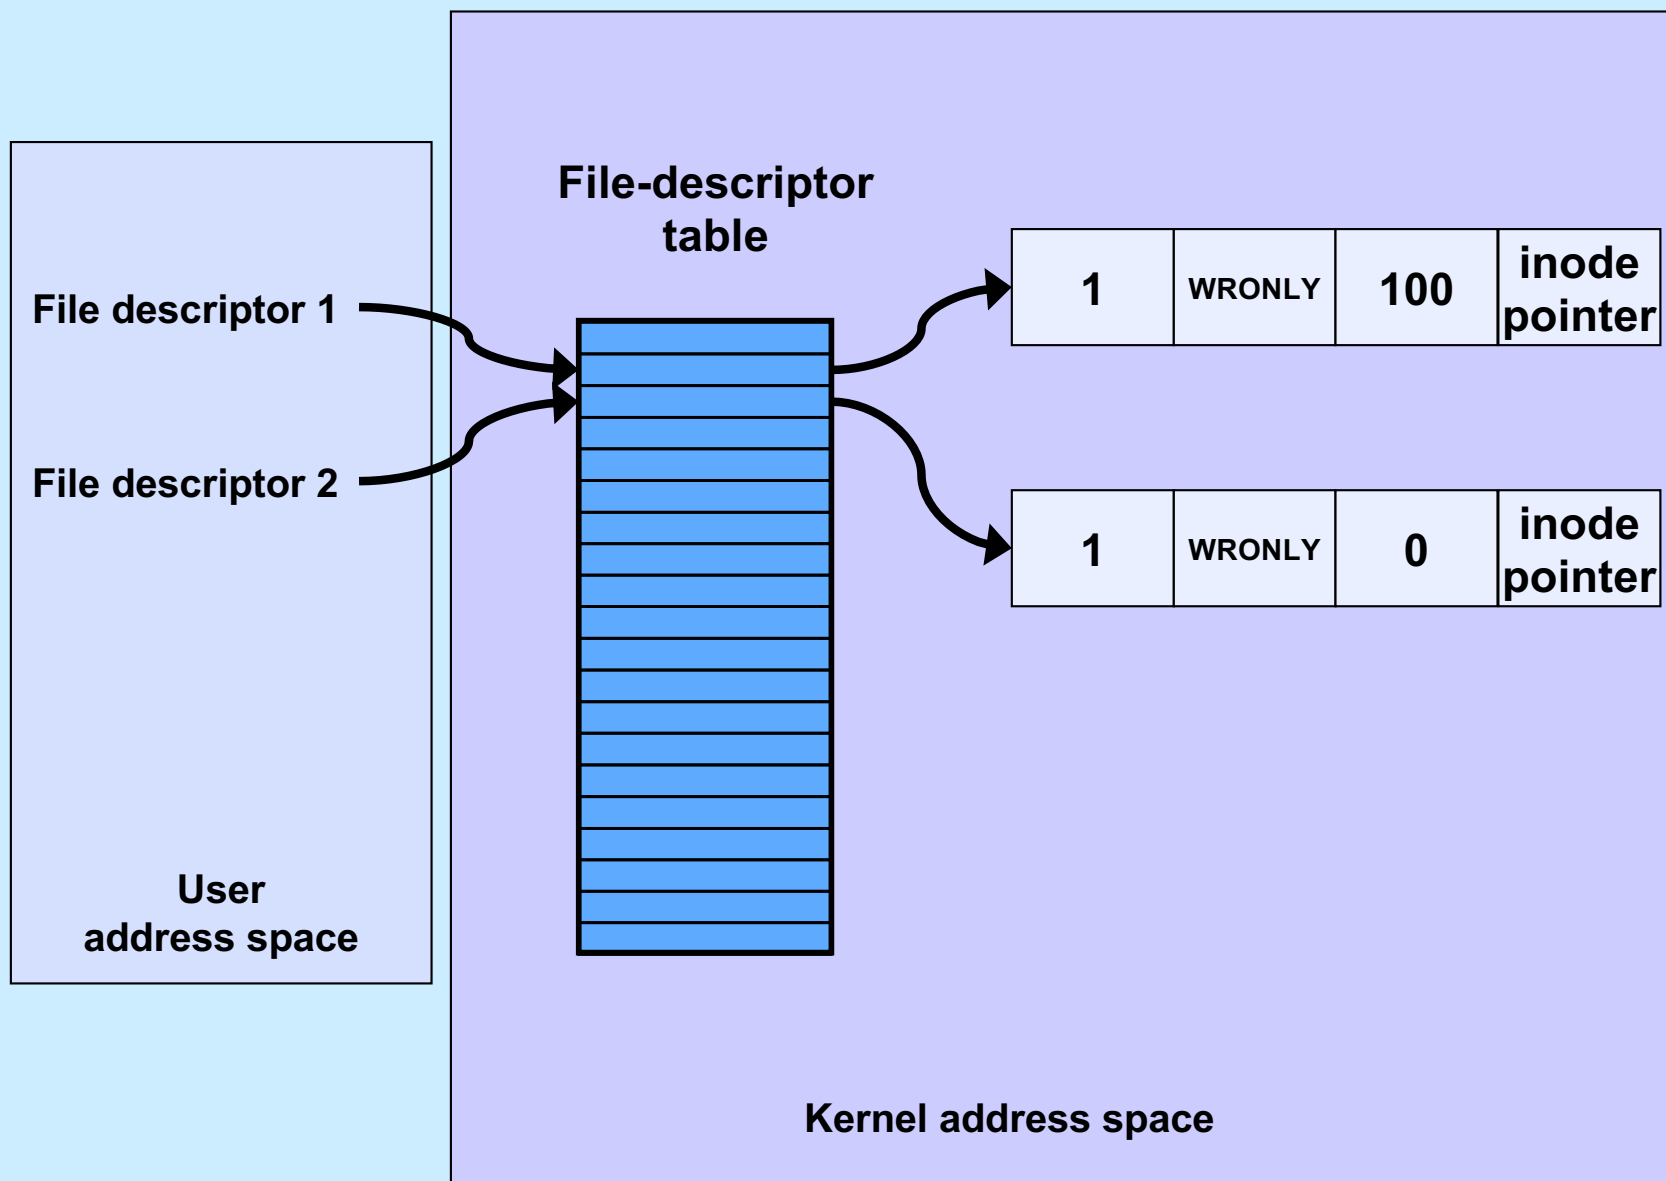
```
$ echon 1 >Output 2>Output
```

– **both stdout and stderr go to Output file**

# Redirected Output



# Redirected Output After Write



# Quiz 1

- **Suppose we run**

```
$ echo 3 >Output 2>Output
```

- **The input line is**

X

- **What is the final content of Output?**

a) X\nX\nreps too large, reduced to 2\n

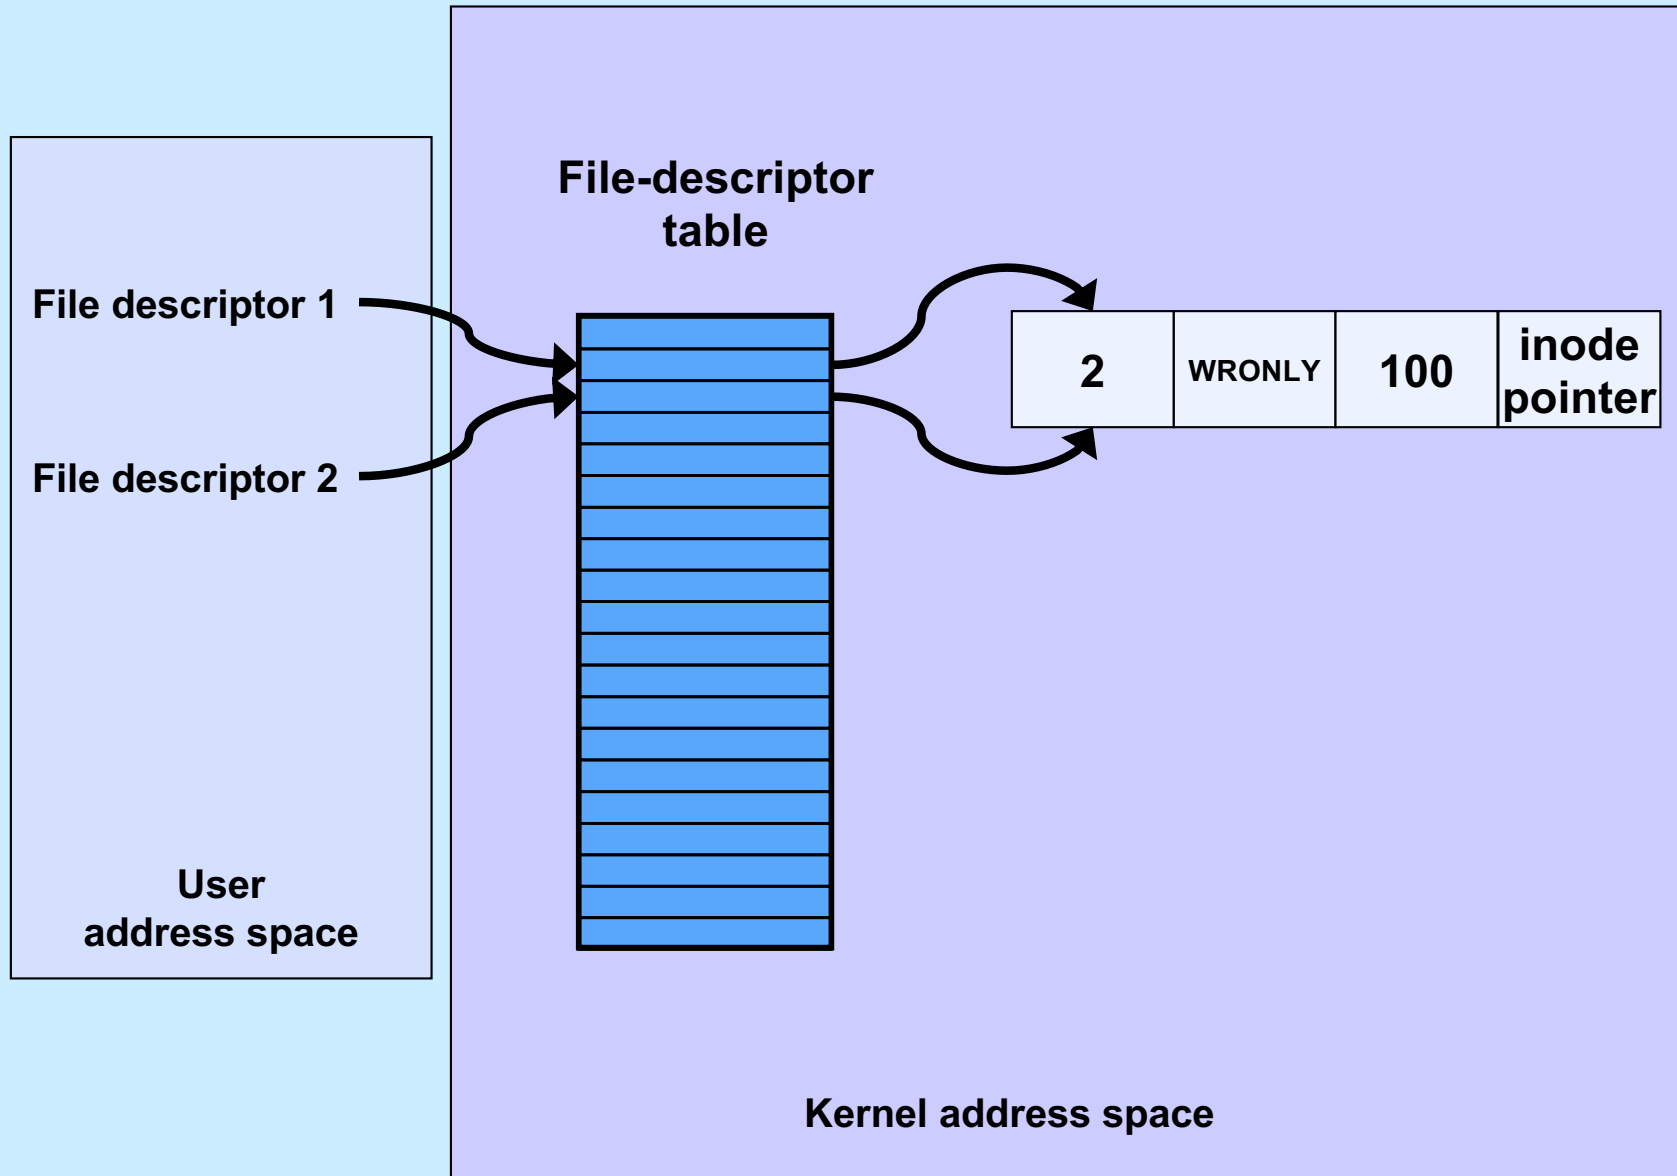
b) X\nX\n too large, reduced to 2\n

c) reps too large, reduced to 2\nX\nX\n

# Sharing Context Information

```
if (fork() == 0) {  
    /* set up file descriptors 1 and 2 in the child process */  
    close(1);  
    close(2);  
    if (open("/home/twd/Output", O_WRONLY) == -1) {  
        exit(1);  
    }  
    dup(1); /* set up file descriptor 2 as a duplicate of 1 */  
    char *argv[] = {"echon", 2};  
    execv("/home/twd/bin/echon", argv);  
    exit(1);  
}  
/* parent continues here */
```

# Redirected Output After Dup



# From the Shell ...

```
$ echon 3 >Output 2>&1
```

- **stdout goes to Output file, stderr is the dup of fd 1**

- **with input “X\n” it now produces in Output:**

```
reps too large, reduced to 2\nX\nX\n
```

# Fork and File Descriptors

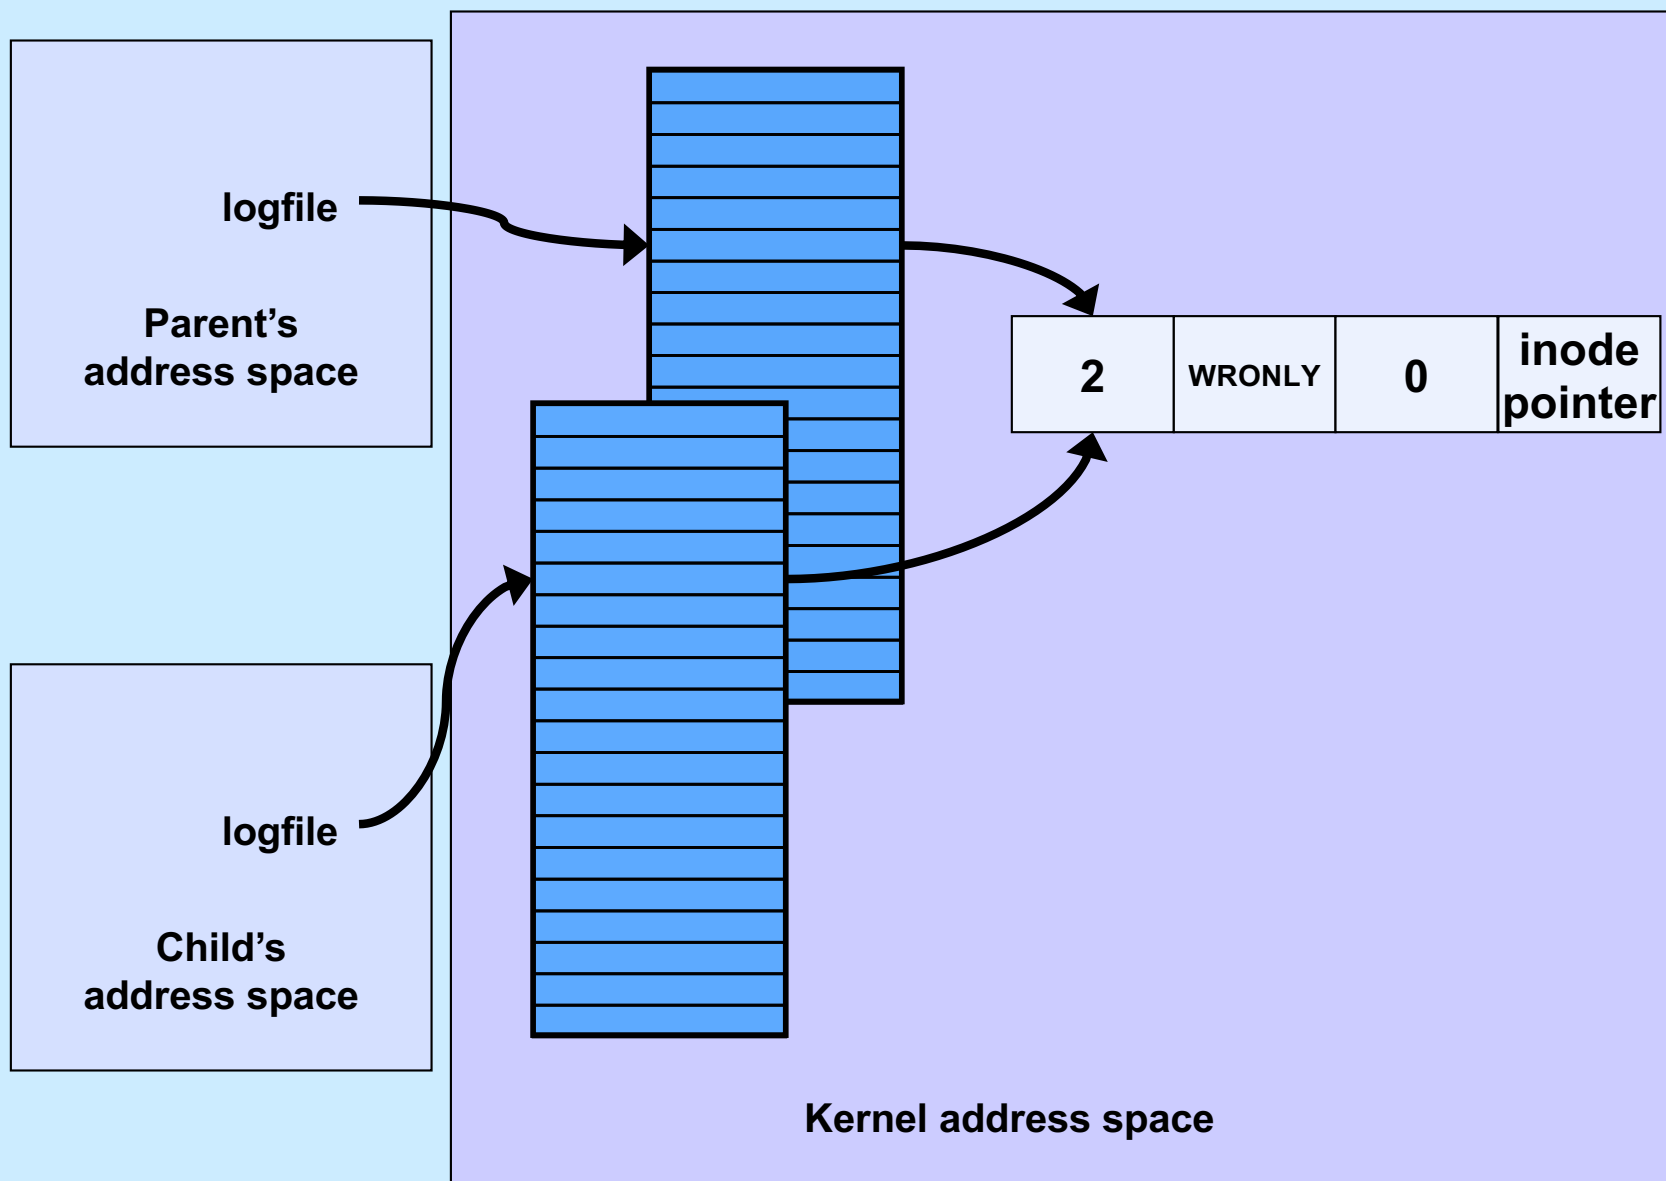
```
int logfile = open("log", O_WRONLY);
if (fork() == 0) {
    /* child process computes something, then does: */
    write(logfile, LogEntry, strlen(LogEntry));
    ...
    exit(0);
}

/* parent process computes something, then does: */

write(logfile, LogEntry, strlen(LogEntry));
...
```



# File Descriptors After Fork



# Quiz 2

```
int main() {  
    if (fork() == 0) {  
        fprintf(stderr, "Child");  
        exit(0);  
    }  
    fprintf(stderr, "Parent");  
}
```


**Suppose the program is run as:**

```
$ prog >file 2>&1
```

**What is the final content of file? (Assume writes are “atomic”.)**

- a) either “Childt” or “Parent”**
- b) either “Child” or “Parent”**
- c) either “ChildParent” or “ParentChild”**

# Directories




# Directory Representation

Component Name	Inode Number
----------------	--------------

directory entry

.	1
..	1
unix	117
etc	4
home	18
pro	36
dev	93

# Hard Links



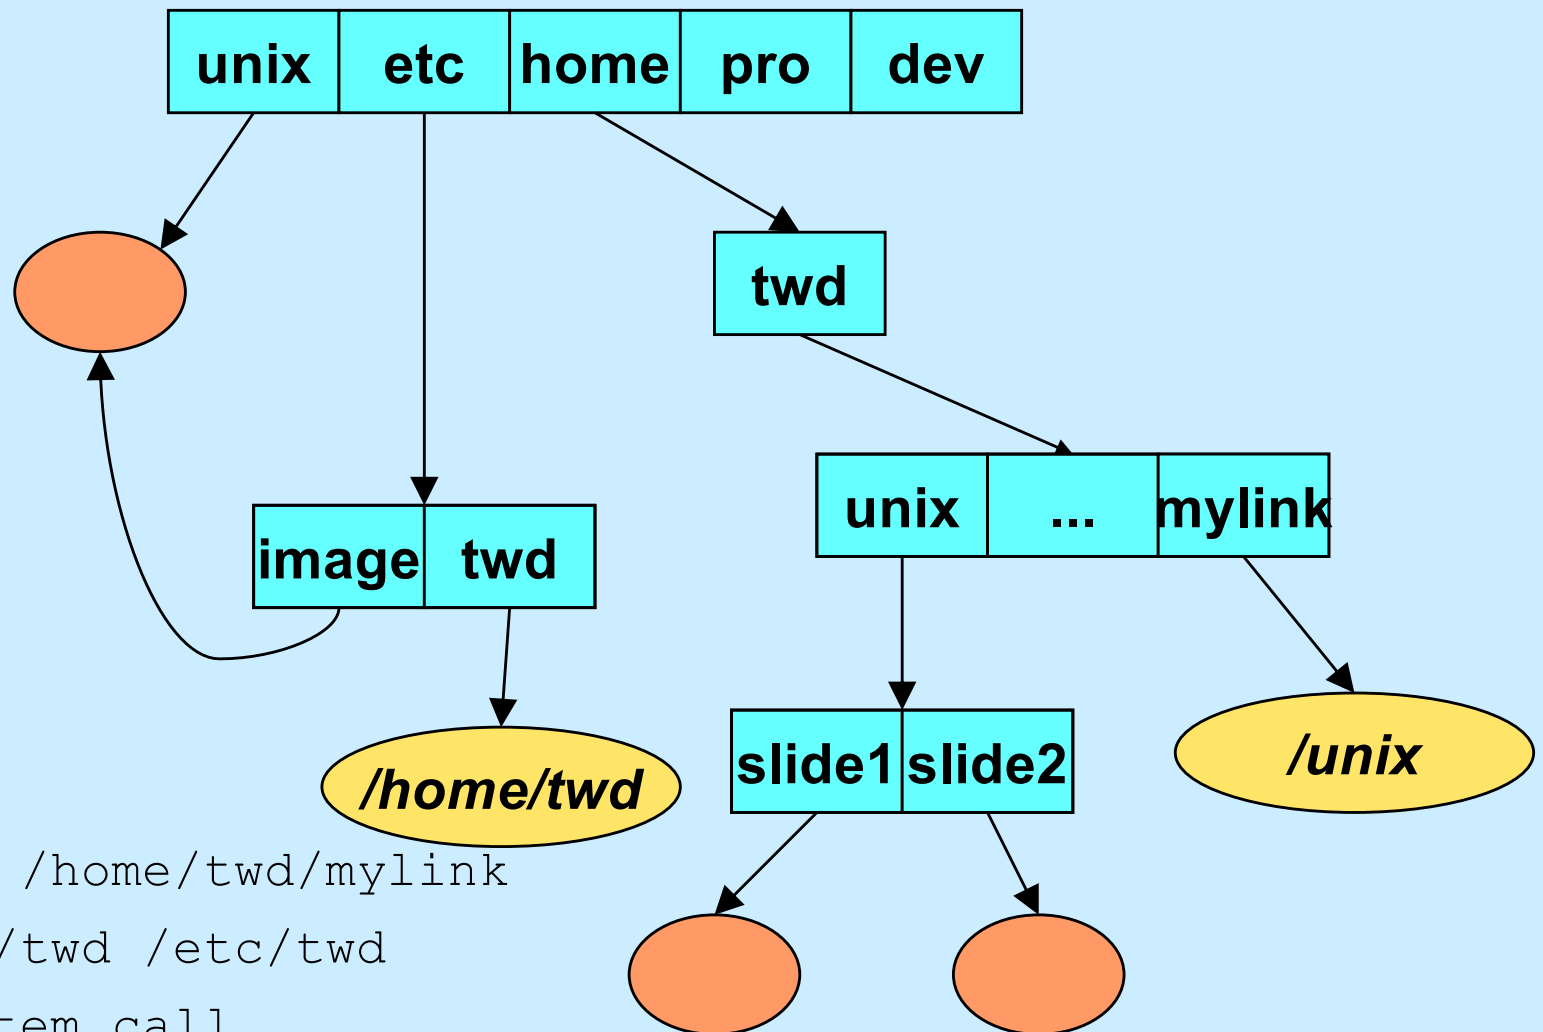
```
$ ln /unix /etc/image  
# link system call
```

# Directory Representation

.	1
..	1
unix	117
etc	4
home	18
pro	36
dev	93

.	4
..	1
image	117
motd	33

# Symbolic Links



```
% ln -s /unix /home/twd/mylink
% ln -s /home/twd /etc/twd
# symlink system call
```

# Working Directory

- **Maintained in kernel for each process**
  - paths not starting from “/” start with the working directory
  - changed by use of the *chdir* system call
    - » *cd* shell command
  - displayed (via shell) using “pwd”
    - » how is this done?



# Open

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int open(const char *path, int options [, mode_t mode])
```

## – options

- » **O\_RDONLY**      open for reading only
- » **O\_WRONLY**      open for writing only
- » **O\_RDWR**      open for reading and writing
- » **O\_APPEND**      set the file offset to *end of file* prior to each *write*
- » **O\_CREAT**      if the file does not exist, then create it, setting its mode to *mode* adjusted by *umask*
- » **O\_EXCL**      if **O\_EXCL** and **O\_CREAT** are set, then *open* fails if the file exists
- » **O\_TRUNC**      delete any previous contents of the file

# Appending Data to a File (1)

```
int fd = open("file", O_WRONLY);  
lseek(fd, 0, SEEK_END);  
    // sets the file location to the end  
write(fd, buffer, bsize);  
    // does this always write to the  
    // end of the file?
```

# Appending Data to a File (2)

```
int fd = open("file", O_WRONLY | O_APPEND);  
write(fd, buffer, bsize);  
    // this is guaranteed to write to the  
    // end of the file
```

# In the Shell ...

**% program >> file**

# File Access Permissions

- **Who's allowed to do what?**
  - **who**
    - » **user (owner)**
    - » **group**
    - » **others (rest of the world)**
  - **what**
    - » **read**
    - » **write**
    - » **execute**

# Permissions Example

**adm group:  
joe, angie**

```
$ ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 joe      adm      1024 Dec 17 13:34 A
```

```
drwxr----- 2 joe      adm      1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-   1 joe      adm        593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```

```
-r--rw-rw-   1 joe      adm        446 Dec 17 13:34 x
```

```
-rw----rw-   1 angie    adm        446 Dec 17 13:45 y
```

# Setting File Permissions

```
#include <sys/types.h>
#include <sys/stat.h>
int chmod(const char *path, mode_t mode)
```

- sets the file permissions of the given file to those specified in *mode*
- only the owner of a file and the superuser may change its permissions
- nine combinable possibilities for *mode* (*read/write/execute for user, group, and others*)
  - » S\_IRUSR (0400), S\_IWUSR (0200), S\_IXUSR (0100)
  - » S\_IRGRP (040), S\_IWGRP (020), S\_IXGRP (010)
  - » S\_IROTH (04), S\_IWOTH (02), S\_IXOTH (01)

# Permission Bits

- **It's worth your while to remember this!**
  - read: 4
  - write: 2
  - execute: 1
  - read/write: 6
  - read/write/execute: 7
- **user:group:others**
  - » **0751**
    - rwx for user, rx for group, x for others
  - » **0640**
    - rw for user, r for group, nothing for others



# Umask

- **Standard programs create files with “maximum needed permissions” as mode**
  - **compilers: 0777**
  - **editors: 0666**
- **Per-process parameter, *umask*, used to turn off undesired permission bits**
  - **e.g., turn off all permissions for others, write permission for group: set umask to 027**
    - » **compilers: permissions =  $0777 \& \sim(027) = 0750$**
    - » **editors: permissions =  $0666 \& \sim(027) = 0640$**
  - **set with *umask* system call or (usually) shell command**

# Quiz 3

**You get the following message when you attempt to execute `./program` (a file that you own):**

**`bash: ./program: Permission denied`**

**Your first response should be:**

- a) execute the shell command  
`chmod 0644 program`**
- b) execute the shell command  
`chmod 0755 program`**
- c) find the source code for program and recompile it**
- d) make an Ed post**

# Creating a File

- Use either *open* or *creat*
  - `open(const char *pathname, int flags, mode_t mode)`
    - » flags must include `O_CREAT`
  - `creat(const char *pathname, mode_t mode)`
    - » `open` is preferred
- The *mode* parameter helps specify the permissions of the newly created file
  - `permissions = mode & ~umask`